

It's all text

A model only ever reads **text**. Whatever you hand it — a PDF, a spreadsheet, a web page — is first flattened to plain text, chopped into **tokens**, turned into numbers, and read **in order**. It writes tokens back, in order. **Text in, text out.**

STRUCTURED FREE-FORM

.txt	Plain text No formatting at all — just characters. <i>lecture transcripts, oral-history interviews, primary sources, field notes, Shakespeare's sonnets</i>
.rtf	Rich text A little formatting: bold, italics, headings. <i>handouts, recommendation letters, annotated readings</i>
.md	Markdown Light formatting from a few simple marks. The lingua franca of this work. <i>syllabi, lecture notes, reading guides, problem sets, rubrics, prompts & skills, CLAUDE.md</i>
.html	Web markup Tagged structure a browser can render. <i>course sites, online readings, digital exhibits — like this handout</i>
.csv	Comma-separated values Rows and columns — a table in plain text. <i>gradebooks, survey results, coded interview data, corpus & archive metadata, lab measurements</i>
.json	Structured data Labelled key / value pairs. Strict rules. <i>citation records, structured annotations, dataset & config files, API responses</i>
.py .js	Code Scripts that actually <i>run</i> . The strictest rules of all. <i>text-mining & NLP, data cleaning, statistics, visualizations, grading automations</i>

200K – 1M tokens →

